

# Clustering Assignments

- ❖ Assignment done by :- Dev Mulchandani

- ❖ Part-C :- Gaussian mixture models clustering

In Part C, I performed clustering using Gaussian Mixture Models, which assume that the data is generated from a mixture of several Gaussian distributions. After preprocessing and scaling the numerical features, I trained GMM models with different numbers of components and selected the best model using the Bayesian Information Criterion (BIC). I then predicted the cluster memberships based on the probability distributions learned by the model. To assess clustering quality, I calculated the silhouette score and visualized the resulting clusters in PCA-reduced two dimensions. This approach showed how GMM can model more complex cluster shapes compared to K-Means.

- ❖ Screenshots:-

The screenshot shows a Jupyter Notebook interface. The title bar displays the notebook's name: "part\_c\_gaussian\_mixture\_models(Dev\_M).ipynb". Below the title bar is a menu bar with options: File, Edit, View, Insert, Runtime, Tools, Help. Underneath the menu bar is a toolbar with icons for Commands, Code, Text, and Run all. The main content area has a sidebar on the left with icons for file operations (New, Open, Save, etc.) and a section titled "Part (c) – Gaussian Mixture Models (GMM)". The main content pane contains the following text:

```
Steps:  
1. Load dataset (default: Wine Quality).  
2. Standardize numeric features.  
3. Fit GMMs with different component counts and compare BIC.  
4. Choose best model, compute silhouette score, and visualize.
```

At the bottom of the content pane, there is a box containing the text: "Assignment done by :- Dev Mulchandani".

## Load dataset (upload / URL / Kaggle)

```
[2] ✓ 11s # @title Load dataset (upload / URL / Kaggle)
import pandas as pd
import zipfile
from pathlib import Path

try:
    from google.colab import files # type: ignore
    IN_COLAB = True
except Exception:
    IN_COLAB = False

DEFAULT_URL = "https://raw.githubusercontent.com/riorain/wine-quality-dataset/master/winequality-red.csv" # You can change this to another CSV URL

print("How do you want to load the dataset?")
print("1 = upload CSV file manually")
print("2 = download from URL (uses DEFAULT_URL above)")
print("3 = download from Kaggle (you must provide kaggle.json & dataset name)")
choice = input("Enter 1, 2, or 3: ").strip()

if choice == "1":
    if not IN_COLAB:
        raise RuntimeError("Manual upload only works in Google Colab.")
    uploaded = files.upload()
    fname = list(uploaded.keys())[0]
    df = pd.read_csv(fname)
    print("Loaded:", fname, "shape:", df.shape)
elif choice == "2":
    if not DEFAULT_URL:
        raise ValueError("DEFAULT_URL is empty. Please set it to a valid CSV URL or choose another option.")
    df = pd.read_csv(DEFAULT_URL)
    print("Loaded from URL. Shape:", df.shape)
elif choice == "3":
    import os, subprocess

    if IN_COLAB:
        from google.colab import files # type: ignore
        print("Please upload your kaggle.json file (from your Kaggle account).")
        uploaded = files.upload()
        kaggle_path = Path("~/kaggle").expanduser()
        kaggle_path.mkdir(parents=True, exist_ok=True)
        for fn in uploaded:
            Path(fn).replace(kaggle_path / "kaggle.json")
        os.chmod(kaggle_path / "kaggle.json", 0o600)

    # Install Kaggle CLI
    import sys
    !pip -q install kaggle

DATASET_SLUG = input("Enter Kaggle dataset slug (e.g. 'uciml/iris'): ").strip()

# Download entire dataset (may contain multiple files)
!kaggle datasets download -d $DATASET_SLUG -p kaggle_data
```

```

# Unzip everything
kaggle_dir = Path("kaggle_data")
kaggle_dir.mkdir(exist_ok=True)
for zpath in kaggle_dir.glob("*.zip"):
    with zipfile.ZipFile(zpath, "r") as zf:
        zf.extractall(kaggle_dir)

csv_files = list(kaggle_dir.rglob("*.csv"))
if not csv_files:
    raise FileNotFoundError("No CSV files found in Kaggle dataset; please inspect kaggle_data/ manually.")
csv_path = csv_files[0]
print("Using CSV:", csv_path)
df = pd.read_csv(csv_path)
print("Loaded from Kaggle. Shape:", df.shape)
else:
    raise ValueError("Invalid choice. Please run this cell again.")

df.head()

```

... How do you want to load the dataset?  
 1 = upload CSV file manually  
 2 = download from URL (uses DEFAULT\_URL above)  
 3 = download from Kaggle (you must provide kaggle.json & dataset name)

Enter 1, 2, or 3: 1

winequality.csv

winequality.csv(text/csv) - 18817 bytes, last modified: 02/12/2025 - 100% done

Saving winequality.csv to winequality.csv

Loaded: winequality.csv shape: (200, 5)

	fixed_acidity	volatile_acidity	citric_acid	residual_sugar	alcohol	
0	8.170528	0.332532	0.605048	6.349688	13.079245	
1	10.992366	1.434196	0.847980	1.100501	10.305749	
2	9.162480	0.116799	0.457128	6.576316	10.000033	
3	6.286952	1.418771	0.551276	7.615411	12.197512	
4	5.963014	1.284253	0.741329	2.603784	12.374327	

Next steps: [Generate code with df](#) [New interactive sheet](#)

## Preprocess data

```

[3]
✓ 1s
# @title Preprocess data
import numpy as np
from sklearn.preprocessing import StandardScaler

X = df.select_dtypes(include=["float64", "int64"]).values
print("Feature matrix shape:", X.shape)

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

```

Feature matrix shape: (200, 5)

✓ Fit Gaussian Mixture models and select best BIC

---

[4] ✓ 1s

▶ # @title Fit Gaussian Mixture models and select best BIC  
from sklearn.mixture import GaussianMixture  
import numpy as np  
  
lowest\_bic = np.inf  
best\_gmm = None  
best\_k = None  
bics = []  
  
for k in range(2, 8):  
 gmm = GaussianMixture(n\_components=k, covariance\_type="full", random\_state=42)  
 gmm.fit(X\_scaled)  
 bic = gmm.bic(X\_scaled)  
 bics.append(bic)  
 print(f"k={k}, BIC={bic:.2f}")  
 if bic < lowest\_bic:  
 lowest\_bic = bic  
 best\_gmm = gmm  
 best\_k = k  
  
print("Best k according to BIC:", best\_k)  
labels = best\_gmm.predict(X\_scaled)

✓ ...  
k=2, BIC=3029.47  
k=3, BIC=3017.01  
k=4, BIC=3117.27  
k=5, BIC=3197.99  
k=6, BIC=3248.42  
k=7, BIC=3329.47  
Best k according to BIC: 3

---

## Evaluate & visualize

```
[5] ✓ 0s
▶ # @title Evaluate & visualize
from sklearn.metrics import silhouette_score
from sklearn.decomposition import PCA
import matplotlib.pyplot as plt

score = silhouette_score(X_scaled, labels)
print(f"Silhouette score (k={best_k}): {score:.3f}")

pca = PCA(n_components=2)
X_2d = pca.fit_transform(X_scaled)

plt.figure()
plt.scatter(X_2d[:, 0], X_2d[:, 1], c=labels)
plt.title("GMM Clusters (PCA projection)")
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.show()
```

... Silhouette score (k=3): 0.129

