


Clustering Assignments

❖ Assignment done by :- Dev Mulchandani

❖ Part-G :- Document Clustering Using LLM Embeddings

In Part G, I performed document clustering using modern large-language-model embeddings. I first loaded a collection of text documents and converted each one into a dense vector using a pretrained SentenceTransformer (MiniLM) model. These embeddings captured the semantic meaning of the documents, allowing similar topics to cluster together more effectively than with traditional bag-of-words methods. After applying K-Means to the embeddings, I evaluated the clusters and inspected representative documents from each group. This showed how LLM embeddings can produce high-quality text clusters in NLP applications.

❖ Screenshots:-



The screenshot shows a Jupyter Notebook interface with the following details:

- Title Bar:** Shows the file name "part_g_document_clustering_embeddings(Dev_M).ipynb" and standard menu options: File, Edit, View, Insert, Runtime, Tools, Help.
- Toolbar:** Includes search, code, text, and run all buttons.
- Table of Contents:** On the left, there is a sidebar with icons for files, cells, and help. Below it, a section titled "Part (g) – Document Clustering using LLM Embeddings" is expanded, showing a list of steps.
- Content Area:** The main area displays the steps for the assignment:
 1. Load a text dataset (default: 20 Newsgroups sample).
 2. Encode documents using a SentenceTransformer model.
 3. Cluster embeddings with K-Means.
 4. Inspect sample documents from each cluster.
- Footer:** A box at the bottom contains the text "Assignment done by :- Dev Mulchandani".

✓ Load dataset (upload / URL / Kaggle)

```
[1] ✓ t2s
# @title Load dataset (upload / URL / Kaggle)
import pandas as pd
import zipfile
from pathlib import Path

try:
    from google.colab import files # type: ignore
    IN_COLAB = True
except Exception:
    IN_COLAB = False

DEFAULT_URL = "https://raw.githubusercontent.com/selva86/datasets/master/newsgroups.csv" # You can change t

print("How do you want to load the dataset?")
print("1 = upload CSV file manually")
print("2 = download from URL (uses DEFAULT_URL above)")
print("3 = download from Kaggle (you must provide kaggle.json & dataset name)")
choice = input("Enter 1, 2, or 3: ").strip()

if choice == "1":
    if not IN_COLAB:
        raise RuntimeError("Manual upload only works in Google Colab.")
    uploaded = files.upload()
    fname = list(uploaded.keys())[0]
    df = pd.read_csv(fname)
    print("Loaded:", fname, "shape:", df.shape)
elif choice == "2":
    if not DEFAULT_URL:
        raise ValueError("DEFAULT_URL is empty. Please set it to a valid CSV URL or choose another option.")
    df = pd.read_csv(DEFAULT_URL)
    print("Loaded from URL. Shape:", df.shape)
elif choice == "3":
    import os, subprocess

    if IN_COLAB:
        from google.colab import files # type: ignore
        print("Please upload your kaggle.json file (from your Kaggle account).")
        uploaded = files.upload()
        kaggle_path = Path("~/kaggle").expanduser()
        kaggle_path.mkdir(parents=True, exist_ok=True)
        for fn in uploaded:
            Path(fn).replace(kaggle_path / "kaggle.json")
            os.chmod(kaggle_path / "kaggle.json", 0o600)

    # Install kaggle CLI
    import sys
    !pip -q install kaggle
```

```

DATASET_SLUG = input("Enter Kaggle dataset slug (e.g. 'uciml/iris'): ").strip()

# Download entire dataset (may contain multiple files)
!kaggle datasets download -d $DATASET_SLUG -p kaggle_data

# Unzip everything
kaggle_dir = Path("kaggle_data")
kaggle_dir.mkdir(exist_ok=True)
for zpath in kaggle_dir.glob("*.zip"):
    with zipfile.ZipFile(zpath, "r") as zf:
        zf.extractall(kaggle_dir)

csv_files = list(kaggle_dir.rglob("*.csv"))
if not csv_files:
    raise FileNotFoundError("No CSV files found in Kaggle dataset; please inspect kaggle_data/ manually.")
csv_path = csv_files[0]
print("Using CSV:", csv_path)
df = pd.read_csv(csv_path)
print("Loaded from Kaggle. Shape:", df.shape)
else:
    raise ValueError("Invalid choice. Please run this cell again.")

df.head()

```

... How do you want to load the dataset?
 1 = upload CSV file manually
 2 = download from URL (uses DEFAULT_URL above)
 3 = download from Kaggle (you must provide kaggle.json & dataset name)
 Enter 1, 2, or 3: 1

documents.csv
documents.csv(text/csv) - 3585 bytes, last modified: 02/12/2025 - 100% done
 Saving documents.csv to documents.csv
 Loaded: documents.csv shape: (100, 1)

	text	grid
0	The stock market rallied today.	
1	Deep learning improves AI capabilities.	
2	Football is a popular sport.	
3	Space exploration expands our knowledge.	
4	Healthy eating leads to better life.	

Next steps: [Generate code with df](#) [New interactive sheet](#)

▼ Install sentence-transformers and extras

[2] ✓ 34s

```

# @title Install sentence-transformers and extras
!pip -q install sentence-transformers umap-learn hdbscan

import numpy as np
from sentence_transformers import SentenceTransformer

```

▼ Encode documents

```
[3] ✓ 6s
▶ # @title Encode documents
text_col_candidates = [c for c in df.columns if c.lower() in ["text", "article", "content", "message"]]
if text_col_candidates:
    text_col = text_col_candidates[0]
else:
    text_col = df.columns[0]
print("Using text column:", text_col)

documents = df[text_col].astype(str).tolist()
max_docs = 3000
documents = documents[:max_docs]

model = SentenceTransformer("all-MiniLM-L6-v2")
embeddings = model.encode(documents, batch_size=64, show_progress_bar=True)
embeddings.shape
```

```
... Using text column: text
/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens).
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
    warnings.warn(
modules.json: 100% [██████████] 349/349 [00:00<00:00, 16.0kB/s]
config_sentence_transformers.json: 100% [██████████] 116/116 [00:00<00:00, 8.91kB/s]
README.md: [██] 10.5k/? [00:00<00:00, 198kB/s]
sentence_bert_config.json: 100% [██████████] 53.0/53.0 [00:00<00:00, 1.20kB/s]
config.json: 100% [██████████] 612/612 [00:00<00:00, 22.7kB/s]
model.safetensors: 100% [██████████] 90.9M/90.9M [00:01<00:00, 67.1MB/s]
tokenizer_config.json: 100% [██████████] 350/350 [00:00<00:00, 7.38kB/s]
vocab.txt: [██] 232k/? [00:00<00:00, 2.37MB/s]
tokenizer.json: [██] 466k/? [00:00<00:00, 8.15MB/s]
special_tokens_map.json: 100% [██████████] 112/112 [00:00<00:00, 6.93kB/s]
config.json: 100% [██████████] 190/190 [00:00<00:00, 3.40kB/s]
Batches: 100% [██████████] 2/2 [00:01<00:00, 1.19it/s]
(100, 384)
```

▼ Cluster embeddings

```
[4] ✓ 1s # @title Cluster embeddings
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
k = 10
kmeans = KMeans(n_clusters=k)
labels = kmeans.fit_predict(embeddings)

score = silhouette_score(embeddings, labels)
print(f"Silhouette score: {score:.3f}")

...
... Silhouette score: 1.000
/usr/local/lib/python3.12/dist-packages/sklearn/base.py:1389: (
    return fit_method(estimator, *args, **kwargs)
```

✓ Inspect clusters with example texts

```
[5] ✓ 0s
# @title Inspect clusters with example texts
from collections import defaultdict

cluster_examples = defaultdict(list)
for idx, label in enumerate(labels):
    if len(cluster_examples[label]) < 5:
        cluster_examples[label].append(documents[idx])

for label, examples in cluster_examples.items():
    print("-" * 80)
    print(f"Cluster {label}")
    print("-" * 80)
    for ex in examples:
        print("-", ex[:200].replace("\n", " ") + "...")
    print()
```

*** ====== Cluster 6 ======

- The stock market rallied today....

===== Cluster 4 =====

- Deep learning improves AI capabilities....

===== Cluster 0 =====

- Football is a popular sport....

===== Cluster 1 =====

- Space exploration expands our knowledge....

===== Cluster 3 =====

- Healthy eating leads to better life....

===== Cluster 2 =====

- The stock market rallied today....
- The stock market rallied today....
- The stock market rallied today....
- The stock market rallied today....