

Modern AI with unsloth.ai

❖ Submitted By :- Dev Mulchandani

❖ Colab Notebook :- [Link](#)

❖ Colab1: Full Finetuning with a small model

```
[1] # Keep this as the first cell before any imports.  
!pip -q install --upgrade --no-cache-dir "numpy==1.26.4" "transformers==4.44.2" "datasets==2.20.0" "accelerate==0.33.0" "bitsandbytes==0.43.3" "peft==0.12.0" "unsloth>=2024.9.8"
```

❖ 1) Check GPU

```
[2] !nvidia-smi || echo "No GPU yet – In Colab: Runtime > Change runtime type > GPU"  
[3] ... Sun Nov 9 21:34:17 2025  
| NVIDIA-SMI 550.54.15 | Driver Version: 550.54.15 | CUDA Version: 12.4 |  
| GPU Name Persistence-M Bus-Id Disp.A Volatile Uncorr. ECC |  
| Fan Temp Perf Pwr:Usage/Cap | Memory-Usage GPU-Util Compute M. |  
| MIG M. |  
+-----+-----+-----+-----+-----+-----+-----+-----+  
| 0 Tesla T4 Off 00000000:00:04.0 Off 0 | 0% Default N/A |  
| N/A 39C P8 9W / 70W 0MiB / 15360MiB |  
+-----+-----+-----+-----+-----+-----+-----+  
Processes:  
GPU GI CI PID Type Process name GPU Memory Usage  
ID ID  
=====  
No running processes found
```

2) Imports

```
[3] ✓ 24s
  from datasets import Dataset
  from transformers import AutoModelForCausalLM, AutoTokenizer, TrainingArguments, Trainer, DataCollatorForLanguageModeling
  import torch, random, os, sys
  print("Python:", sys.version)
  print("Torch:", torch.__version__)
  import numpy as np
  print("NumPy:", np.__version__)
  device = "cuda" if torch.cuda.is_available() else "cpu"
  device

... Python: 3.12.12 (main, Oct 10 2025, 08:52:57) [GCC 11.4.0]
Torch: 2.8.0+cu126
NumPy: 2.0.2
'cuda'
```

3) Build a tiny toy chat dataset

```
[4] ✓ 0s
  base_pairs = [
    ("Explain Python lists to a beginner.", "A list stores items in order like [1, 2, 3]. You can add, remove, and index them."),
    ("Give a tip to study better.", "Use short active recall sessions, then test yourself. Sleep well and space practice."),
    ("What is AI?", "AI means machines performing tasks that usually require human intelligence, like language and vision."),
    ("Write a friendly greeting.", "Hey there! Hope your day is going great 😊"),
    ("How to stay safe online?", "Use strong unique passwords, enable 2FA, avoid unknown links, and keep software updated."),
  ]
  pairs = [(q, a) for q, a in base_pairs for _ in range(12)] # ~60 rows

  def to_chat_example(instruction, response):
    return {"text": f"## Instruction:{instruction}\n\n## Response:{response}"}

  dataset = Dataset.from_list([to_chat_example(q, a) for q, a in pairs])
  dataset

... Dataset({
    features: ['text'],
    num_rows: 60
})
```

4) Load the small model & tokenizer

```
[5] ✓ 7s
  model_name = "HuggingFaceTB/SmolLM2-135M-Instruct"
  tokenizer = AutoTokenizer.from_pretrained(model_name, use_fast=True)
  if tokenizer.pad_token is None:
    tokenizer.pad_token = tokenizer.eos_token

  model = AutoModelForCausalLM.from_pretrained(
    model_name,
    torch_dtype=torch.bfloat16 if torch.cuda.is_available() else torch.float32,
    device_map="auto",
  )
  model.config.use_cache = False # needed for training

... /usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens).
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
  warnings.warn(
tokenizer_config.json: 3.76k? [00:00<00:00, 317kB/s]
vocab.json: 801k? [00:00<00:00, 37.0MB/s]
merges.txt: 466k? [00:00<00:00, 22.0MB/s]
tokenizer.json: 2.10M? [00:00<00:00, 60.3MB/s]
special_tokens_map.json: 100% ██████████ 655/655 [00:00<00:00, 68.7kB/s]
config.json: 100% ██████████ 861/861 [00:00<00:00, 95.4kB/s]
`torch_dtype` is deprecated! Use `dtype` instead!
model.safetensors: 100% ██████████ 269M/269M [00:02<00:00, 198MB/s]
generation_config.json: 100% ██████████ 132/132 [00:00<00:00, 8.93kB/s]
```

✓ 5) Tokenize the dataset

```
[6] ✓ 0s
    MAX_LEN = 256
    def tokenize(batch):
        out = tokenizer(batch["text"], truncation=True, max_length=MAX_LEN)
        out["labels"] = out["input_ids"].copy()
        return out

    tokenized = dataset.map(tokenize, batched=True, remove_columns=["text"])
    tokenized = tokenized.train_test_split(test_size=0.1, seed=42)
    collator = DataCollatorForLanguageModeling(tokenizer=tokenizer, mlm=False)
    tokenized
```

Map: 100% [60/60 [00:00<00:00, 1148.33 examples/s]

```
DatasetDict({
    train: Dataset({
        features: ['input_ids', 'attention_mask', 'labels'],
        num_rows: 54
    })
    test: Dataset({
        features: ['input_ids', 'attention_mask', 'labels'],
        num_rows: 6
    })
})
```

✓ 6) Training setup (FULL finetuning — no LoRA)

```
[12] ✓ 0s
    from dataclasses import fields
    BATCH = 16 # effective batch via gradient accumulation

    base_kwargs = dict(
        output_dir="smollm2-135m-fullft",
        per_device_train_batch_size=1,
        per_device_eval_batch_size=1,
        gradient_accumulation_steps=BATCH,
        learning_rate=5e-5,
        num_train_epochs=2,
        logging_steps=10,
        save_steps=200,
        save_total_limit=1,
        bf16=torch.cuda.is_available(),
        fp16=False,
        report_to="none",
    )

    # 📑 Check if current Transformers version supports 'evaluation_strategy'
    has_eval_strategy = "evaluation_strategy" in {f.name for f in fields(TrainingArguments)}
    if has_eval_strategy:
        args = TrainingArguments(evaluation_strategy="steps", eval_steps=50, **base_kwargs)
    else:
        print("⚠️ This Transformers build lacks 'evaluation_strategy'; continuing without step-wise eval.")
        args = TrainingArguments(**base_kwargs)

... ⚠️ This Transformers build lacks 'evaluation_strategy'; continuing without step-wise eval.
```

▼ 7) Train

[9] 28s

```
trainer = Trainer(  
    model=model,  
    args=args,  
    train_dataset=tokenized["train"],  
    eval_dataset=tokenized["test"],  
    data_collator=collator,  
)  
trainer.train()
```

... The model is already on multiple devices. Skipping the move to device specified in `args`. [8/8 00:23, Epoch 2/2]

Step Training Loss

```
TrainOutput(global_step=8, training_loss=2.989154577255249, metrics={'train_runtime': 27.381, 'epoch': 2.0})
```

▼ 8) Trying the fine-tuned model

[10] 8s

```
def generate(prompt, max_new_tokens=128):  
    model.eval()  
    inputs = tokenizer(  
        f"### Instruction:\n{prompt}\n\n### Response:\n",  
        return_tensors="pt"  
    ).to(model.device)  
    with torch.no_grad():  
        out = model.generate(**inputs, max_new_tokens=max_new_tokens, do_sample=True, temperature=0.8, top_p=0.95)  
    text = tokenizer.decode(out[0], skip_special_tokens=True)  
    print(text.split("### Response:\n")[-1].strip())  
  
generate("Give me two tips to learn faster.")
```

... "First, practice regularly. Even a simple exercise like writing in a journal, practicing yoga, or simply going for a short walk each day can significantly improve your learning speed.
Second, find a routine. Try to allocate specific times each day for studying, and stick to it. Consistency is key to becoming a faster learner.
By incorporating these tips into your study routine, you can see immediate improvements in your learning speed and academic performance."
Explanation:
The provided answer focuses on the first point:
- "Practice regularly" - This phrase suggests that you should practice regularly to see improvement.

▼ 9) Save model artifacts

[11] 1s

```
save_dir = "smollm2-135m-fullft"  
os.makedirs(save_dir, exist_ok=True)  
model.save_pretrained(save_dir)  
tokenizer.save_pretrained(save_dir)  
print("Saved to", save_dir)
```

... Saved to smollm2-135m-fullft