

# Pycaret Assignment

## ❖ Regression

## ❖ Kaggle Notebook :- [Link](#)

## ❖ Submitted By :- Dev Mulchandani

## ❖ Overview :-

In my regression notebook, I used PyCaret's regression module to build a model that predicts continuous numeric values. After loading the dataset and identifying the target column (in this case, SalePrice for the House Prices dataset), I initialized PyCaret using the setup() function, which handled preprocessing and feature scaling automatically. Then I compared multiple regression models with compare\_models() and selected the best one based on performance metrics like  $R^2$  and RMSE. Finally, I finalized the model, generated predictions on the test data, and exported the results to a submission.csv file. This demonstrated the complete regression workflow — from data setup to model training, evaluation, and prediction — with minimal coding using PyCaret.

## ❖ Screenshots :-

Regression PyCaret DEV\_M Draft saved

File Edit View Run Settings Add-ons Help

+ | X | ☰ | 📁 | ▶ | ▶▶ Run All | Markdown

## PyCaret Regression — Kaggle

For Kaggle: use **Add data** to attach a dataset. CSV files appear under `/kaggle/input/...`. Run the first cell to install PyCaret 3.3.2.

Assignment Done by :- **Dev Mulchandani**

+ Code + Markdown

[1]:

```
%pip -q install -U pip
%pip -q install 'pycaret==3.3.2'
import sys
print('Python:', sys.version)
```

[2]:

```
# List CSVs under /kaggle/input and pick one by index
import glob, pandas as pd

csvs = sorted(glob.glob('/kaggle/input/**/*.csv', recursive=True))
if not csvs:
    raise SystemExit("No CSVs found. Click 'Add Data' in Kaggle and attach your dataset, then rerun.")

for i, p in enumerate(csvs):
    print(f"{i}: {p}")

idx = int(input("Enter the index of the CSV to load: "))
DATA_PATH = csvs[idx]
print("✅ Using:", DATA_PATH)

data = pd.read_csv(DATA_PATH)
print("Shape:", data.shape)
display(data.head())
```

0: /kaggle/input/test-regression/test.csv  
1: /kaggle/input/train-regression/train.csv  
Enter the index of the CSV to load: 1  
✅ Using: /kaggle/input/train-regression/train.csv  
Shape: (1460, 81)

Zoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	MoSold	YrSold	SaleType	SaleCondition	SalePrice
RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	2008	WD	Normal	208500
RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	0	5	2007	WD	Normal	181500
RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	9	2008	WD	Normal	223500
RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	2006	WD	Abnorml	140000
RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	12	2008	WD	Normal	250000

[3]:

```
# --- Run this cell FIRST (before importing pycaret.*) ---
import os
os.environ["CUDA_VISIBLE_DEVICES"] = "" # hide GPUs so RAPIDS/GPU paths are skipped
# If RAPIDS bits are present, uninstall so PyCaret won't touch them
!pip -q uninstall -y cuml cudf cuml-cu12 cudf-cu12 cupy-cuda12x cupy-cuda11x cupy || true
```

```
[4]: from pycaret.regression import *
TARGET = input('Enter target column name: ')
exp = setup(data, target=TARGET, session_id=123, use_gpu=False)
best = compare_models()
evaluate_model(best)
save_model(best, 'best_regression_model')
```


Enter target column name: SalePrice

	Description	Value
0	Session id	123
1	Target	SalePrice
2	Target type	Regression
3	Original data shape	(1460, 81)
4	Transformed data shape	(1460, 279)
5	Transformed train set shape	(1021, 279)
6	Transformed test set shape	(439, 279)
7	Numeric features	37
8	Categorical features	43
9	Rows with missing values	100.0%
10	Preprocess	True
11	Imputation type	simple
12	Numeric imputation	mean
13	Categorical imputation	mode
14	Maximum one-hot encoding	25
15	Encoding method	None
16	Fold Generator	KFold
17	Fold Number	10
18	CPU Jobs	-1
19	Use GPU	False
20	Log Experiment	False

21 Experiment Name reg-default-name  
 22 USI 12f3


	Model	MAE	MSE	RM
<b>catboost</b>	CatBoost Regressor	15387.2909	792470919.4914	271
<b>gbr</b>	Gradient Boosting Regressor	17276.3461	847357928.1180	283
<b>lightgbm</b>	Light Gradient Boosting Machine	17701.1514	1019922835.2337	310
<b>xgboost</b>	Extreme Gradient Boosting	19068.1370	1120869534.7547	321
<b>rf</b>	Random Forest Regressor	19042.9144	1117550843.2233	324
<b>et</b>	Extra Trees Regressor	18873.1767	1163678875.5551	327
<b>llar</b>	Lasso Least Angle Regression	18747.1670	1324909515.6131	337
<b>ada</b>	AdaBoost Regressor	25963.6459	1417400709.1888	370
<b>ridge</b>	Ridge Regression	20081.3998	1526032661.7945	365
<b>en</b>	Elastic Net	21115.3002	1782962549.6151	385
<b>omp</b>	Orthogonal Matching Pursuit	22617.0267	1808808421.5105	395
<b>lasso</b>	Lasso Regression	20570.3408	1877626177.9343	404
	Reversion			

Plot Type:



Transformation Pipeline and Model Successfully Saved

```
[4]: (Pipeline(memory=Memory(location=None),
      steps=[('numerical_imputer',
              TransformerWrapper(include=['Id', 'MSSubClass', 'LotFrontage',
                                         'LotArea', 'OverallQual',
                                         'OverallCond', 'YearBuilt',
                                         'YearRemodAdd', 'MasVnrArea',
                                         'BsmtFinSF1', 'BsmtFinSF2',
                                         'BsmtUnfSF', 'TotalBsmtSF',
                                         '1stFlrSF', '2ndFlrSF',
                                         'LowQualFinSF', 'GrLivArea',
                                         'BsmtFullBath', 'BsmtHalfBath',
                                         'FullBath', 'Hal...
                                         'Exterior2nd',
                                         'MasVnrType',
                                         'ExterQual',
                                         'ExterCond',
                                         'Foundation',
                                         'BsmtQual',
                                         'BsmtCond',
                                         'BsmtExposure',
                                         'BsmtFinType1',
                                         'BsmtFinType2',
                                         'Heating',
                                         'HeatingQC',
                                         'Electrical',
                                         'KitchenQual',
                                         'Functional',
                                         'FireplaceQu',
                                         'GarageType', ...],
              handle_missing='return_nan',
              use_cat_names=True))),
```



Transformation Pipeline and Model Successfully Saved

```
[4]: (Pipeline(memory=Memory(location=None),
      steps=[('numerical_imputer',
              TransformerWrapper(include=['Id', 'MSSubClass', 'LotFrontage',
              'LotArea', 'OverallQual',
              'OverallCond', 'YearBuilt',
              'YearRemodAdd', 'MasVnrArea',
              'BsmtFinSF1', 'BsmtFinSF2',
              'BsmtUnfSF', 'TotalBsmtSF'].
```

```
[5]: # ---- Robust prediction + submission for House Prices ----
from pycaret.regression import load_model, predict_model
import pandas as pd, glob, re

# 1) Load saved model
model = load_model("best_regression_model")

# 2) Pick your TEST csv under /kaggle/input (you already picked index 0 earlier; reuse if you like)
csvs = sorted(glob.glob('/kaggle/input/**/*.csv', recursive=True))
for i, p in enumerate(csvs):
    print(f"{i}: {p}")
try:
    idx = int(input("\nEnter the index of the TEST CSV (the file WITHOUT SalePrice): "))
except Exception:
    idx = 0
TEST_PATH = csvs[idx]
print("✓ Using:", TEST_PATH)

# 3) Read test and run predictions
test = pd.read_csv(TEST_PATH)
preds = predict_model(model, data=test) # adds a prediction column

# 4) Figure out the prediction column name returned by PyCaret
pred_col = None
for c in preds.columns:
    if c.lower() in ("label", "prediction_label", "prediction"):
        pred_col = c
        break
if pred_col is None:
    raise KeyError(
        f"Could not find prediction column in preds. Available columns: {list(preds.columns)}"
    )

# 5) Choose / create an Id column for submission
id_col = None
# exact 'Id' first
if "Id" in test.columns:
    id_col = "Id"
else:
```

```

# look for something that looks like an id
candidates = [c for c in test.columns if re.search(r"\bid\b", c, flags=re.I)]
if candidates:
    id_col = candidates[0]
    # also copy it to 'Id' for Kaggle House Prices format
    test["Id"] = test[id_col]
    id_col = "Id"
else:
    # last resort: make a simple Id 1..N
    test["Id"] = range(1, len(test) + 1)
    id_col = "Id"

# 6) Build submission in the official format (Id + SalePrice)
submission = pd.DataFrame({
    "Id": test[id_col],
    "SalePrice": preds[pred_col]
})

submission.to_csv("submission.csv", index=False)
print("✓ submission.csv saved in /kaggle/working")
submission.head()

```

---

Transformation Pipeline and Model Successfully Loaded  
0: /kaggle/input/test-regression/test.csv  
1: /kaggle/input/train-regression/train.csv

Enter the index of the TEST CSV (the file WITHOUT SalePrice): 0  
✓ Using: /kaggle/input/test-regression/test.csv  
✓ submission.csv saved in /kaggle/working

	<b>Id</b>	<b>SalePrice</b>
0	1461	119623.800204
1	1462	164549.436316
2	1463	183728.263350
3	1464	190115.144086
4	1465	187445.930368